

DAT565/DIT407 Assignment 6

Ola Bratt
ola.bratt@gmail.com

Patrick Attimont
patrickattimont@gmail.com

2024-02-xx

This paper is addressing the assignment 6 study queries within the *Introduction to Data Science & AI* course, DIT407 at the University of Gothenburg and DAT565 at Chalmers. The main source of information for this project is derived from the lectures and Skiena [1]. Assignment 6 is about neural networks.

Problem 1: The dataset

The dataset consists of 60,000 training images and 10,000 test images. Each image is a 28x28 pixel grayscale image. The images are labeled with the corresponding digit. A random set of the images are shown in figure 1.




Figure 1: MNIST images

Problem 2: Single hidden layer

The neural network with a single hidden layer has 784 (28×28) input nodes, 300 hidden nodes, and 10 output nodes. The activation function is ReLU for the hidden layer with a batch normalization and logarithmic softmax for the output layer.

```
Epoch [1/10], Training Loss: 0.2200, Test Loss: 0.1065, Test Accuracy: 0.9695
Epoch [2/10], Training Loss: 0.1057, Test Loss: 0.0921, Test Accuracy: 0.9725
Epoch [3/10], Training Loss: 0.0768, Test Loss: 0.0728, Test Accuracy: 0.9765
Epoch [4/10], Training Loss: 0.0613, Test Loss: 0.0666, Test Accuracy: 0.9793
Epoch [5/10], Training Loss: 0.0497, Test Loss: 0.0643, Test Accuracy: 0.9791
Epoch [6/10], Training Loss: 0.0421, Test Loss: 0.0639, Test Accuracy: 0.9798
```

Epoch [7/10], Training Loss: 0.0355, Test Loss: 0.0628, Test Accuracy: 0.9811
Epoch [8/10], Training Loss: 0.0309, Test Loss: 0.0569, Test Accuracy: 0.9823
Epoch [9/10], Training Loss: 0.0279, Test Loss: 0.0616, Test Accuracy: 0.9796
Epoch [10/10], Training Loss: 0.0225, Test Loss: 0.0590, Test Accuracy: 0.9816

Accuracy for single hidden layer: 0.9816




Figure 2: Single hidden layer

Problem 3: Two hidden layers

Accuracy for two hidden layers: 0.9855




Figure 3: Two hidden layers

Problem 4: Convolutional neural network

Accuracy for convolutional neural network: 0.9929




Figure 4: Convolutional neural network

References

- [1] Steven S Skiena. *The Data Science Design Manual*. Retrieved 2024-01-20. 2024. URL: <https://ebookcentral.proquest.com/lib/gu/detail.action?docID=6312797>.

Appendix: Source Code

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 import torch.optim as optim
5 import torchvision.transforms as transforms
6 import matplotlib.pyplot as plt
7 import pandas as pd
8 import numpy as np
9 import matplotlib.pyplot as plt
10
11 from torchvision import datasets
12 from torch.utils.data import DataLoader
13
14
15 class NeuralNet(nn.Module):
16     def __init__(self, input_size, hidden_sizes, output_size):
17         super(NeuralNet, self).__init__()
18         layer_sizes = [input_size] + hidden_sizes + [output_size]
19
20         layers = []
21         for i in range(len(layer_sizes) - 1):
22             layers.append(nn.Linear(layer_sizes[i], layer_sizes[i
23                                     ↪ +1]))
24             if i < len(layer_sizes) - 2: # Add ReLU and batch
25                                     ↪ normalization except for the last layer
26                 layers.append(nn.BatchNorm1d(layer_sizes[i+1]))
27                 layers.append(nn.ReLU())
28             else:
29                 layers.append(nn.LogSoftmax(dim=1)) #layers.append(
30                                     ↪ nn.Softmax(dim=1))
31
32         self.model = nn.Sequential(*layers)
33
34     def forward(self, x):
35         x = x.view(-1, 28 * 28)
36         x = self.model(x)
37         return x
38
39 class CNN(nn.Module):
40     def __init__(self):
41         super(CNN, self).__init__()
42         self.conv1 = nn.Conv2d(in_channels=1, out_channels=16,
43                                 ↪ kernel_size=3, stride=1, padding=1)
44         self.conv2 = nn.Conv2d(in_channels=16, out_channels=32,
45                                 ↪ kernel_size=3, stride=1, padding=1)
46         self.fc1 = nn.Linear(32 * 7 * 7, 128)
47         self.fc2 = nn.Linear(128, 10)
48
49     def forward(self, x):
50         x = F.relu(self.conv1(x))
51         x = F.max_pool2d(x, kernel_size=2, stride=2)
```

```

48         x = F.relu(self.conv2(x))
49         x = F.max_pool2d(x, kernel_size=2, stride=2)
50         x = x.view(-1, 32 * 7 * 7)
51         x = F.relu(self.fc1(x))
52         x = self.fc2(x)
53         x = F.log_softmax(x, dim=1)
54         return x
55
56     def plot_images(dataloader, classes):
57
58         for images, labels in train_loader:
59             print("Image-shape:", images.size())
60             print("Label-shape:", labels.size())
61
62             fig = plt.figure(figsize=(10, 10))
63             for i in range(4):
64                 plt.subplot(5, 5, i + 1)
65                 plt.imshow(images[i].squeeze(), cmap='gray')
66                 plt.title(f'Label: {labels[i]}')
67                 plt.axis('off')
68             plt.show()
69             fig.savefig('mnist_images.png', bbox_inches='tight')
70             break
71
72     def calculate_accuracy(model, dataloader):
73         correct = 0
74         total = 0
75         model.eval()
76
77         with torch.no_grad():
78             for inputs, labels in dataloader:
79                 outputs = model(inputs)
80                 _, predicted = torch.max(outputs, 1)
81                 correct += (predicted == labels).sum().item()
82                 total += labels.size(0)
83
84         accuracy = correct / total
85
86         return accuracy
87
88
89     def train(model, criterion, optimizer, train_loader, test_loader,
90             ↪ num_epochs, name):
91         train_losses = []
92         test_losses = []
93
94         for epoch in range(num_epochs):
95             model.train()
96             running_loss = 0.0
97
98             for images, labels in train_loader:
99                 outputs = model(images)
100                 loss = criterion(outputs, labels)
101
102                 optimizer.zero_grad()
103                 loss.backward()
104                 optimizer.step()
105
106                 running_loss += loss.item()
107
108         epoch_loss = running_loss / len(train_loader)

```

```

109         train_losses.append(epoch_loss)
110
111     model.eval()
112     correct = 0
113     total = 0
114     test_loss = 0.0
115     with torch.no_grad():
116         for images, labels in test_loader:
117             outputs = model(images)
118             _, predicted = torch.max(outputs, 1)
119             correct += (predicted == labels).sum().item()
120             total += labels.size(0)
121             loss = criterion(outputs, labels)
122             test_loss += loss.item()
123     accuracy = correct / total
124     test_loss /= len(test_loader)
125     test_losses.append(test_loss)
126
127     print(f"Epoch- [{epoch+1}/{num_epochs}], Training-Loss: {
        ↪ epoch_loss:.4f}, Test-Loss: {test_loss:.4f}, Test-
        ↪ Accuracy: {accuracy:.4f}")
128
129
130     fig, ax = plt.subplots(figsize=(8, 6), layout='constrained')
131     ax.plot(range(1, num_epochs + 1), train_losses, label='Training
        ↪ Loss')
132     ax.plot(range(1, num_epochs + 1), test_losses, label='Test-Loss
        ↪ ')
133     ax.set_xlabel('Epochs')
134     ax.set_ylabel('Loss')
135     ax.set_title('Training and Test Losses')
136     ax.legend()
137     plt.show()
138     fig.savefig(name + ".pdf", bbox_inches='tight')
139
140
141
142
143     # Importing the dataset
144     batch_size = 32
145     transform = transforms.Compose([
146         transforms.Resize((28, 28)),
147         transforms.ToTensor(),
148         transforms.Normalize((0.5, ), (0.5,))
149     ])
150
151     train_dataset = datasets.MNIST(root='Assignment6/', train=True,
        ↪ download=True, transform=transform)
152     test_dataset = datasets.MNIST(root='Assignment6/', train=False,
        ↪ download=True, transform=transform)
153
154     train_loader = DataLoader(train_dataset, batch_size=batch_size,
        ↪ shuffle=True, num_workers=2)
155     test_loader = DataLoader(test_dataset, batch_size=batch_size,
        ↪ shuffle=False, num_workers=2)
156
157     print("train-dataset:-", len(train_dataset))
158     print("test-dataset:-", len(test_dataset))
159
160
161     # Single hidden layer
162     input_size = 28 * 28

```

```

163 hidden_sizes = [300]
164 output_size = 10
165
166 modelSHL = NeuralNet(input_size, hidden_sizes, output_size)
167 learning_rate = 0.1
168 optimizer = optim.SGD(modelSHL.parameters(), lr=learning_rate)
169 num_epochs = 10
170 criterion = nn.CrossEntropyLoss()
171
172 train(modelSHL, criterion, optimizer, train_loader, test_loader,
      ↪ num_epochs, "single_hidden_layer")
173 accuracy = calculate_accuracy(modelSHL, test_loader)
174 print(f"Accuracy for single hidden layer: {accuracy:.4f}")
175
176
177 # Two hidden layers
178 hidden_sizes = [500, 300]
179 weight_decay = 0.0001
180 modelTHL = NeuralNet(input_size, hidden_sizes, output_size)
181 optimizer = optim.SGD(modelTHL.parameters(), lr=learning_rate,
      ↪ weight_decay=weight_decay)
182 num_epochs = 40
183
184 train(modelTHL, criterion, optimizer, train_loader, test_loader,
      ↪ num_epochs, "two_hidden_layer")
185 accuracy = calculate_accuracy(modelTHL, test_loader)
186 print(f"Accuracy for two hidden layers: {accuracy:.4f}")
187
188
189 # Convolutional neural network
190 modelCNN = CNN()
191 weight_decay = 0.0001
192 optimizer = optim.SGD(modelCNN.parameters(), lr=learning_rate,
      ↪ weight_decay=weight_decay)
193 num_epochs = 40
194
195 train(modelCNN, criterion, optimizer, train_loader, test_loader,
      ↪ num_epochs, "cnn")
196 accuracy = calculate_accuracy(modelCNN, test_loader)
197 print(f"Accuracy for convolutional neural network: {accuracy:.4f}")

```